

Assignment 3 - Sparse matrix formats on GPUs

Submission Deadline: Mo 29 November, 10am

Task 1: So far we learned about the CSR format. On CPUs this is a widely used standard format. However, it has some severe disadvantages on GPUs, but also on modern vector extensions (AVX, etc.) of CPUs. The paper [Improving the performance of the sparse matrix vector product with GPUs](#) by Vazquez et. al. describes these difficulties and also the alternative Ellpack and Ellpack-R formats that improve on the shortcomings of CSR on GPUs and offer more performance on these devices. Summarize the difficulties of the CSR format and how Ellpack and Ellpack-R solve these difficulties. For a very high mark look for some more modern papers on GPU based matvecs and describe what other modern developments have happened since that paper.

Task 2: Given a sparse matrix in CSR format. Write a CPU based Numba routine that converts this matrix into the Ellpack-R format.

Implement a new class `EllpackMatrix` derived from `scipy.sparse.linalg.LinearOperator`, which in its constructor takes a Scipy sparse matrix in CSR format, converts it to Ellpack-R and provides a routine for matrix-vector product in the Ellpack-R format. At the end the following prototype commands should be possible with your class.

```
my_sparse_mat = EllpackMatrix(csr_mat)
x = numpy.random.randn(my_sparse_mat.shape[1])
y = my_sparse_mat @ x
```

Print to PDF ►

The sparse-matrix vector product at the end shall be executed in the Ellpack-R format.

For your implementation you can either use CPU based Numba code using `prange` for multithreading or write an implementation in Numba-Cuda. For an overall mark of the assignment beyond 80% we would expect a Numba-Cuda implementation (assuming also all other parts are of high standard).

Demonstrate in your solution that your code provides the correct result by verifying for a 1000×1000 sparse random matrix with the standard CSR matvec of sparse matrices and showing for 3 random vectors that the relative distance of your Ellpack-R matvec to the CSR matvec result is in the order of machine precision.

Use the `discretise_poisson` method from the lecture notes to generate the sparse matrix for the Poisson discretization and plot the times for a single matvec for growing matrix-sizes (go as high as you think is reasonable) using the standard Scipy csr matvec and your own Ellpack-R implementation.

Your implementation may not be faster since the matrix derives from a very simple 2d problem with few elements per row. This is not the situation where the additional complexity is usually needed.

Finally, go shopping in the [Matrix Market](#) and try to find two matrices that better show off the Ellpack-R format and do timing comparisons for your chosen matrices.

Previous

◀ [Assignment 2 - GPU Accelerated solution of Poisson problems](#)

Next

[Assignment 4 - Time-dependent problems](#) ▶

By Timo Betcke

© Copyright 2020.