

Assignment4 - Yaru Zeng

December 15, 2021

0.1 Set-up

```
[ ]: # import packages needed for computation

import numpy as np
import matplotlib.pyplot as plt
from scipy.sparse import coo_matrix, identity
from scipy.sparse.linalg import spsolve
import time
from numba import cuda

[ ]: # define a timer for monitoring execution times

class Timer:
    def __enter__(self):
        self.start = time.time()
        return self

    def __exit__(self, *args):
        self.end = time.time()
        self.interval = self.end - self.start
```

0.2 Explicit method: Forward Euler

0.2.1 Implement a finite difference scheme with the explicit time-stepping supported by the Forward Euler method.

For the explicit time stepping I will use the Forward Euler method which gives

$u_{n+1} = u_n + \Delta t f(u_n, t_n)$. In this assignment $f = u_t = \Delta u$, then $u_{n+1} = u_n + \Delta t \Delta u_n$

To solve this problem we consider a discretisation based on the finite difference method. Let $x_i = ih$ with $h = \frac{2}{N-1}$ where $i = 0, \dots, N-1$, then we can find the approximation of $\Delta u_n(i, j)$ by

$$\Delta u_n(i, j) = -[4u_n(i, j) - u_n(i-1, j) - u_n(i+1, j) - u_n(i, j-1) - u_n(i, j+1)]/h^2$$

Finally, the time-stepping problem becomes

$$u_{n+1}(i, j) = u_n(i, j) - C[4u_n(i, j) - u_n(i-1, j) - u_n(i+1, j) - u_n(i, j-1) - u_n(i, j+1)] \text{ where } C = \frac{\Delta t}{h^2}.$$

From the lecture, for the stability of the Forward Euler method for 1-D problem where $u_t = u_{xx}$, then $C = \frac{\Delta t}{h^2} \lesssim \frac{1}{2}$.

In this case which is a 2-D problem, $u_t = u_{xx} + u_{yy}$, the condition becomes $C = \frac{\Delta t}{h^2} \lesssim \frac{1}{4}$.

For simplification, set the time steps as M, then $\Delta t = \frac{1}{M-1}$. With $h = \frac{2}{N-1}$, $C = \frac{\Delta t}{h^2} = \frac{(N-1)^2}{4(M-1)}$.

```
[ ]: # Define a function to execute the Forward Euler method by GPU

@cuda.jit
def forward_euler_gpu(u, u_n1):
    i, j = cuda.grid(2)
    if 0 < i < u.shape[0]-1 and 0 < j < u.shape[1]-1:
        u_n1[i,j] = u[i,j] - C*(4*u[i, j] - u[i-1,j] - u[i+1,j] - u[i,j-1] -
        ↪u[i,j+1])

[ ]: # Execute the Forward Euler method for certain values of N and M to check the
    ↪accuracy

N = 51
M = 10001
t_actual = 0.424011387033

u = np.zeros((N,N), dtype = np.float64)
u[-1,:] = np.float32(5)
u_n1 = np.copy(u)
dt = 1/(M-1)
dx = 2/(N-1)

C = (N-1)**2/(4*(M-1))
u_global_mem = cuda.to_device(u)
u_n1_global_mem = cuda.to_device(u_n1)

SX = 16
SY = 16
nblocks = int(np.ceil(N/SX))

t_star = 0
with Timer() as t:
    while u_global_mem[(N-1)//2,(N-1)//2] < 1.0:
        t_star += dt
        u_n1_global_mem = cuda.to_device(u_n1)
        forward_euler_gpu[(nblocks,nblocks), (SX,SY)](u_global_mem,u_n1_global_mem)
        u_global_mem = u_n1_global_mem

u = u_n1_global_mem.copy_to_host()

rel_err = abs(t_star-t_actual)/abs(t_actual)
```

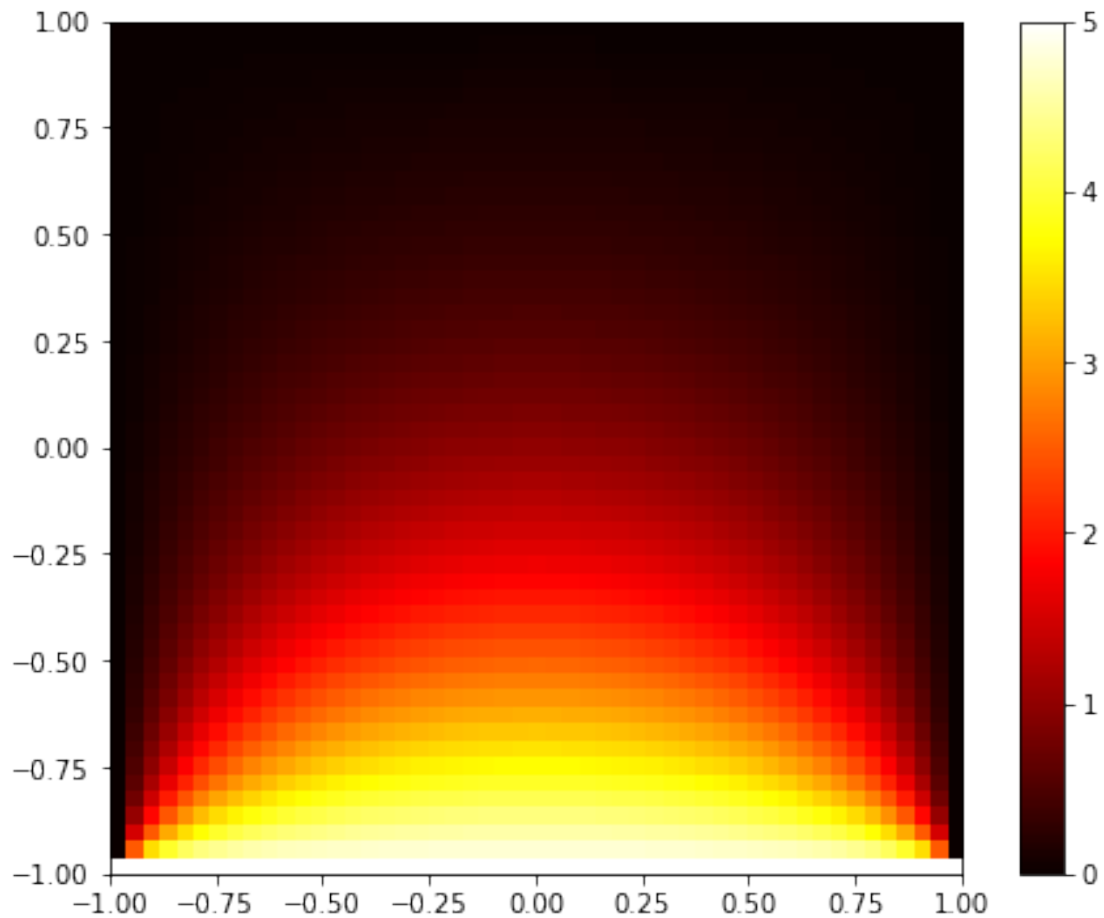
```

print(f"N: {N}, M: {M}, C: {C}")
print(f"t_star: {t_star}")
print(f"t_actual: {t_actual}")
print(f"Relative error: {rel_err}")
print("Time cost: {0} s".format(t.interval))
print("\n")
plt.figure(figsize=(8,6))
plt.imshow(u, extent=(-1, 1, -1, 1), cmap='hot')
plt.colorbar()

```

N: 51, M: 10001, C: 0.0625
 t_star: 0.10590000000000002
 t_actual: 0.424011387033
 Relative error: 0.7502425565949247
 Time cost: 0.6774048805236816 s

[]: <matplotlib.colorbar.Colorbar at 0x7f0d9c0fbf50>



0.2.2 By increasing the number of discretisation points demonstrate how many correct digits you can achieve.

In order to research how the number of discretisation points N and the number of time-steps M influence the number of correct digits, I set a list of C for every certain N provided that when C and N is specified M is a certain number with $C = \frac{\Delta t}{h^2} = \frac{(N-1)^2}{4(M-1)}$.

For the stability of the Forward Euler method, I set C to be less than or equal to 0.25 as mentioned above.

```
[ ]: # set a list of N and C to research how the number of discretisation points N
      ↪and time-steps M influence correct digits

t_actual = 0.424011387033
N_list = [11,101,201,301]
C_list = [0.25,0.125,0.0625,0.03125]

for N in N_list:

    print(f"N = {N}", end="\n")
    print(f"t_actual = {t_actual}")

    t_exe = 0
    table = PrettyTable(["C","M","t_star"],align = 'l')

    for C in C_list:

        @cuda.jit
        def forward_euler_gpu(u, u_n1):
            i, j = cuda.grid(2)
            if 0 < i < u.shape[0]-1 and 0 < j < u.shape[1]-1:
                u_n1[i,j] = u[i,j] - C*(4*u[i, j] - u[i-1,j] - u[i+1,j] - u[i,j-1] -
                ↪u[i,j+1])

        table_data = []
        u = np.zeros((N,N), dtype = np.float64)
        u[0,:] = np.float32(5)
        u_n1 = np.copy(u)
        dx = 2/(N-1)
        dt = C*dx**2
        M = round(1/dt + 1)
        table_data.append(C)
        table_data.append(M)

        u_global_mem = cuda.to_device(u)
        u_n1_global_mem = cuda.to_device(u_n1)
```

```

SX = 16
SY = 16
nblocks = int(np.ceil(N/SX))

t_star = 0
with Timer() as t:
    while u_global_mem[(N-1)//2,(N-1)//2] < 1.0:
        t_star += dt
        u_n1_global_mem = cuda.to_device(u_n1)
        forward_euler_gpu[(nblocks,nblocks),
→(SX,SY)](u_global_mem,u_n1_global_mem)
        u_global_mem = u_n1_global_mem
        t_exe += t.interval

rel_err = abs(t_star-t_actual)/abs(t_actual)
table_data.append(t_star)
table.add_row(table_data)

print(table)
print(f"Time cost: {t_exe} s")
print("\n")

```

```

N = 11
t_actual = 0.424011387033
+-----+-----+-----+
| C      | M    | t_star          |
+-----+-----+-----+
| 0.25   | 101  | 0.4200000000000003 |
| 0.125  | 201  | 0.4200000000000003 |
| 0.0625 | 401  | 0.4225000000000003 |
| 0.03125 | 801  | 0.423749999999999646 |
+-----+-----+-----+
Time cost: 1.2161285877227783 s

```

```

N = 101
t_actual = 0.424011387033
+-----+-----+-----+
| C      | M      | t_star          |
+-----+-----+-----+
| 0.25   | 10001  | 0.423999999999999696 |
| 0.125  | 20001  | 0.4239999999999996146 |
| 0.0625 | 40001  | 0.423999999999999574 |
| 0.03125 | 80001  | 0.4240000000000003417 |
+-----+-----+-----+
Time cost: 43.42120933532715 s

```

```

N = 201
t_actual = 0.424011387033
+-----+-----+-----+
| C      | M      | t_star      |
+-----+-----+-----+
| 0.25   | 40001  | 0.423999999999574 |
| 0.125  | 80001  | 0.4240000000003417 |
| 0.0625 | 160001 | 0.4240062499998455 |
| 0.03125 | 320001 | 0.42400937499942865 |
+-----+-----+-----+
Time cost: 192.81513571739197 s

```

```

N = 301
t_actual = 0.424011387033
+-----+-----+-----+
| C      | M      | t_star      |
+-----+-----+-----+
| 0.25   | 90001  | 0.42400000000018123 |
| 0.125  | 180001 | 0.42401111111106621 |
| 0.0625 | 360001 | 0.4240111111118958 |
| 0.03125 | 720001 | 0.42401111111289236 |
+-----+-----+-----+
Time cost: 533.7136478424072 s

```

From the output, we can find when C is certain, as the number of discretisation points N increases, we get more correct digits. Here the most number of correct digits we get is 6. Furthermore, when N is certain, the number of correct digits will grow as C decreases namely the number of time steps increases. In conclusion, the number of correct digits is relevant to the discretisation points and time steps. When C is certain, the smaller Δt and h is, the more correct digits we can get.

0.2.3 Plot the convergence of the computed time * against the actual time.

In order to observe a better convergence trend, I plot two graphs respectively against N and M . With the same aim, for either of the two graphs, I set the other factor than the x-axis to be a certain value.

For the convergence plot against N , when M is specified as 100001 the maximum value of N is about 317 according to the stability condition $C = \frac{\Delta t}{h^2} \lesssim \frac{1}{4}$.

```

[ ]: # Plot the convergence graph against the number of discretisation points for a
      ↪ certain number of time steps

M = 100001
N_list = [11,31,51,71,91,101,201,301]
err_list = []

```

```

for N in N_list:

    @cuda.jit
    def forward_euler_gpu(u, u_n1):
        i, j = cuda.grid(2)
        if 0 < i < u.shape[0]-1 and 0 < j < u.shape[1]-1:
            u_n1[i,j] = u[i,j] - C*(4*u[i, j] - u[i-1,j] - u[i+1,j] - u[i,j-1] - u
↪u[i,j+1])

    u = np.zeros((N,N), dtype = np.float64)
    u[0,:] = np.float32(5)
    u_n1 = np.copy(u)
    dx = 2/(N-1)
    dt = 1/(M-1)
    C = dt/dx**2

    u_global_mem = cuda.to_device(u)
    u_n1_global_mem = cuda.to_device(u_n1)

    SX = 16
    SY = 16
    nblocks = int(np.ceil(N/SX))

    t_star = 0

    while u_global_mem[(N-1)//2,(N-1)//2] < 1.0:
        t_star += dt
        u_n1_global_mem = cuda.to_device(u_n1)
        forward_euler_gpu[(nblocks,nblocks), (SX,SY)](u_global_mem,u_n1_global_mem)
        u_global_mem = u_n1_global_mem

    t_exe += t.interval

    rel_err = abs(t_star-t_actual)/abs(t_actual)
    err_list.append(rel_err)

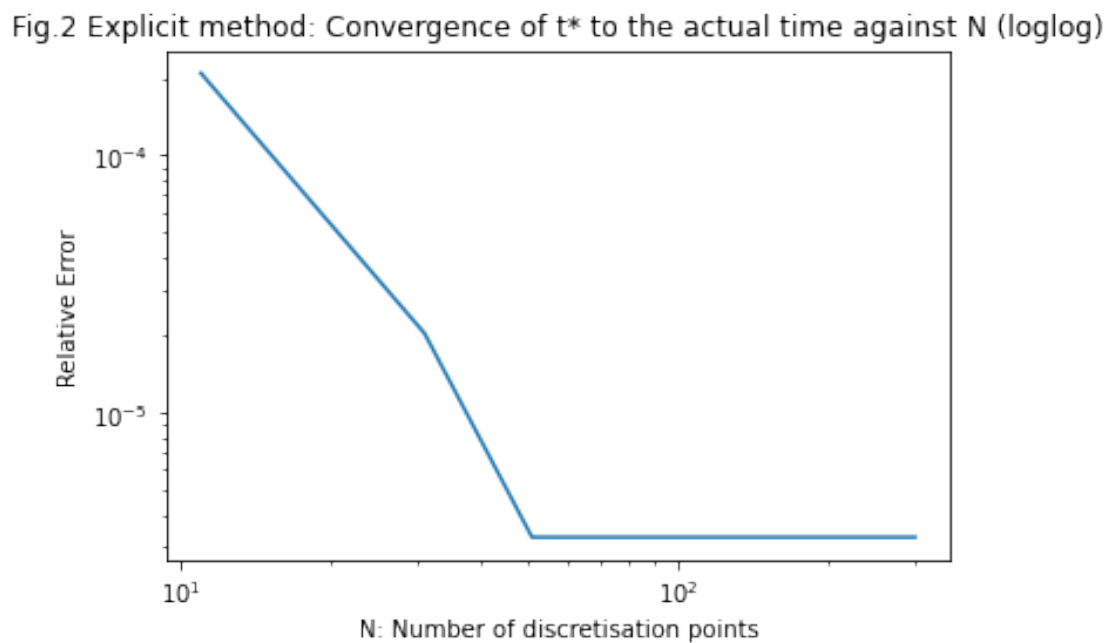
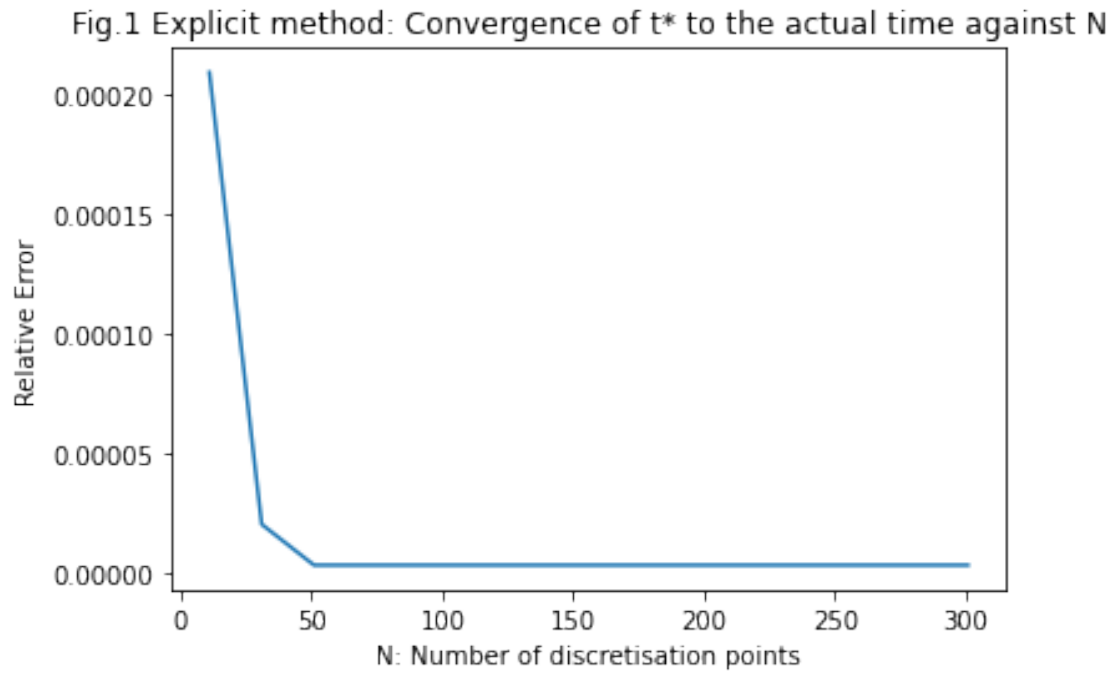
plt.plot(N_list, err_list)
plt.xlabel("N: Number of discretisation points")
plt.ylabel("Relative Error")
plt.title("Fig.1 Explicit method: Convergence of t* to the actual time against_
↪N")
plt.show()

plt.loglog(N_list, err_list)
plt.xlabel("N: Number of discretisation points")

```

```
plt.ylabel("Relative Error")
plt.title("Fig.2 Explicit method: Convergence of  $t^*$  to the actual time against  $N$   

↪  $N$  (loglog)")
plt.show()
```



From the plot, we can see that the convergence of t^* is quadratic. When time steps are certain, more discretisation points will lead to a smaller relative error and a t^* closer to the actual t .

For the convergence plot against M , when N is specified as 51 the minimum value of M is about 2501 according to $C = \frac{\Delta t}{h^2} \lesssim \frac{1}{4}$.

```
[ ]: # Plot the convergence graph against the number of time steps for a certain
      ↪ number of discretisation points

N = 51
M_list = [2501,2801,3001,3501,5001]
err_list = []

for M in M_list:

    @cuda.jit
    def forward_euler_gpu(u, u_n1):
        i, j = cuda.grid(2)
        if 0 < i < u.shape[0]-1 and 0 < j < u.shape[1]-1:
            u_n1[i,j] = u[i,j] - C*(4*u[i, j] - u[i-1,j] - u[i+1,j] - u[i,j-1] -
            ↪ u[i,j+1])

    u = np.zeros((N,N), dtype = np.float64)
    u[0,:] = np.float32(5)
    u_n1 = np.copy(u)
    dx = 2/(N-1)
    dt = 1/(M-1)
    C = dt/dx**2

    u_global_mem = cuda.to_device(u)
    u_n1_global_mem = cuda.to_device(u_n1)

    SX = 16
    SY = 16
    nblocks = int(np.ceil(N/SX))

    t_star = 0
    while u_global_mem[(N-1)//2,(N-1)//2] < 1.0:
        t_star += dt
        u_n1_global_mem = cuda.to_device(u_n1)
        forward_euler_gpu[(nblocks,nblocks), (SX,SY)](u_global_mem,u_n1_global_mem)
        u_global_mem = u_n1_global_mem

    rel_err = abs(t_star-t_actual)/abs(t_actual)
    err_list.append(rel_err)
```

```

plt.plot(M_list, err_list)
plt.xlabel("M: Number of time steps")
plt.ylabel("Relative Error")
plt.title("Fig.3 Explicit method: Convergence of  $t^*$  to the actual time against  $M$ ")
plt.show()

plt.loglog(M_list, err_list)
plt.xlabel("M: Number of time steps")
plt.ylabel("Relative Error")
plt.title("Fig.3 Explicit method: Convergence of  $t^*$  to the actual time against  $M$  (loglog)")
plt.show()

```

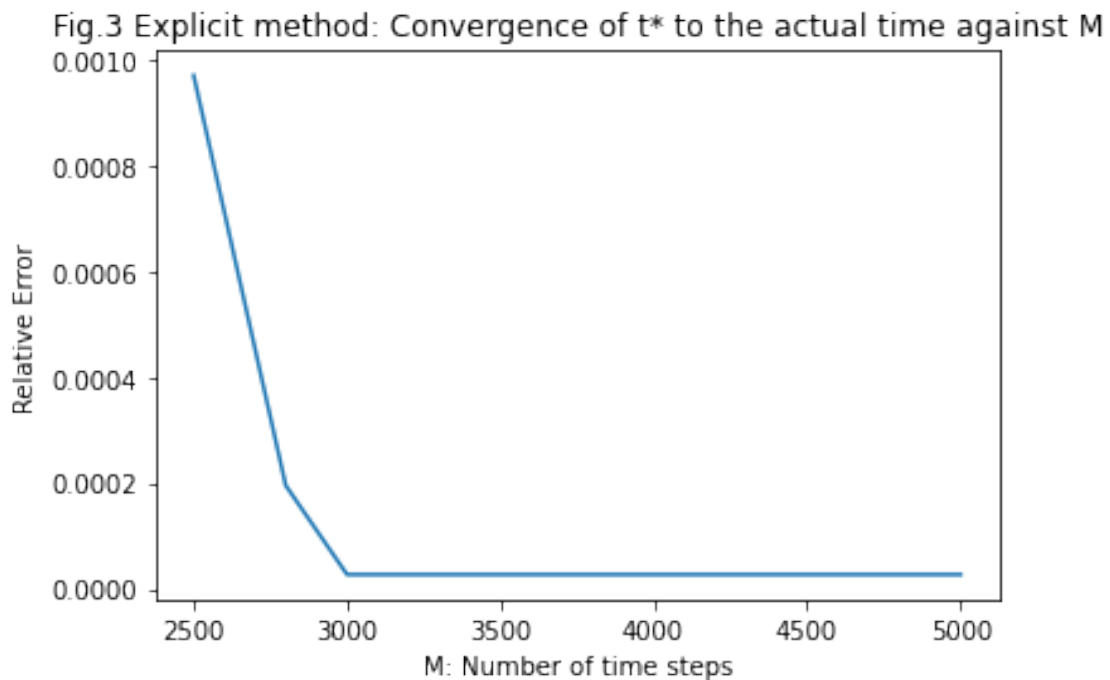
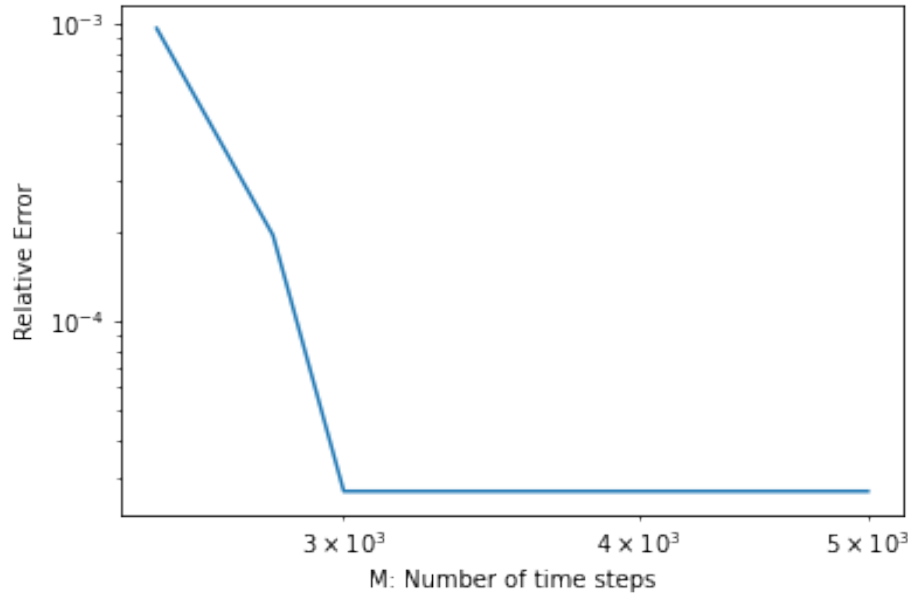


Fig.3 Explicit method: Convergence of t^* to the actual time against M (loglog)



From the graph of convergence of t^* against the number of time steps, we find when discretisation points is certain, we will get a more precise t^* with more time steps.

0.3 Implicit method: Backward Euler

0.3.1 Implement a finite difference scheme with the explicit time-stepping supported by the Backward Euler method.

For the implicit time stepping I will use the Backward Euler method which gives $u_{n+1} = u_n + \Delta t \Delta u_{n+1}$.

To solve this problem, we have to solve a linear equation. For simplifying the algorithm, treat $\Delta u_{n+1} = A u_{n+1}$ where the matrix A is the Laplace operator derived from 5 stencil. Then we get $u_{n+1} = (I - \Delta t A)^{-1} u_n$.

Here, I will use the `discretise_poisson` function on the lecture notes to generate A . To fit the case, I've modified the function by setting $h = \frac{2}{N-1}$ and handling the sign of minus.

```
[ ]: def discretise_poisson(N):
    """Generate the matrix and rhs associated with the discrete Poisson
    ↪ operator."""

    h = 2/(N-1)

    nelements = 5 * N**2 - 16 * N + 16

    row_ind = np.empty(nelements, dtype=np.float64)
    col_ind = np.empty(nelements, dtype=np.float64)
```

```

data = np.empty(nelements, dtype=np.float64)

f = np.empty(N * N, dtype=np.float64)

count = 0
for j in range(N):
    for i in range(N):
        if i == 0 or i == N - 1 or j == 0 or j == N - 1:
            row_ind[count] = col_ind[count] = j * N + i
            data[count] = 1
            f[j * N + i] = 0
            count += 1

        else:
            row_ind[count : count + 5] = j * N + i
            col_ind[count] = j * N + i
            col_ind[count + 1] = j * N + i + 1
            col_ind[count + 2] = j * N + i - 1
            col_ind[count + 3] = (j + 1) * N + i
            col_ind[count + 4] = (j - 1) * N + i

            data[count] = -4 / h**2
            data[count + 1 : count + 5] = 1 / h**2
            f[j * N + i] = 1

            count += 5

return coo_matrix((data, (row_ind, col_ind)), shape=(N**2, N**2)).tocsr(), f

```

```
[ ]: # Define a function to execute the Backward Euler method
```

```

def backward_euler(u, dt):
    u_n = u.reshape((N*N))
    u_n1 = spsolve((I-dt*A),u_n)
    u_n1 = u_n1.reshape((N,N))
    u_n1[-1,:] = 5

    return u_n1

```

```
[ ]: # Execute the Backward Euler method for certain values for N and time_steps M
      →to check the accuracy
```

```

N = 51
M = 10001
t_actual = 0.424011387033

A,_ = discretise_poisson(N)

```

```

I = identity(N*N)

u = np.zeros((N,N), dtype = np.float64)
u[-1,:] = np.float32(5)
dx = 2/(N-1)
C = (N-1)**2/(4*(M-1))
dt = C*dx**2

t_star = 0
with Timer() as t:
    while u[N//2,N//2] < 1:
        u = backward_euler(u,dt)
        t_star += dt

rel_err = abs(t_star-t_actual)/abs(t_actual)
print(f"N: {N}, M: {M}, C: {C}")
print(f"t_star: {t_star}")
print(f"t_actual: {t_actual}")
print(f"Relative error: {rel_err}")
print(f"Time cost: {0} s".format(t.interval))
print("\n")
plt.figure(figsize=(8,6))
plt.imshow(u, extent=(-1, 1, -1, 1), cmap='hot')
plt.colorbar()

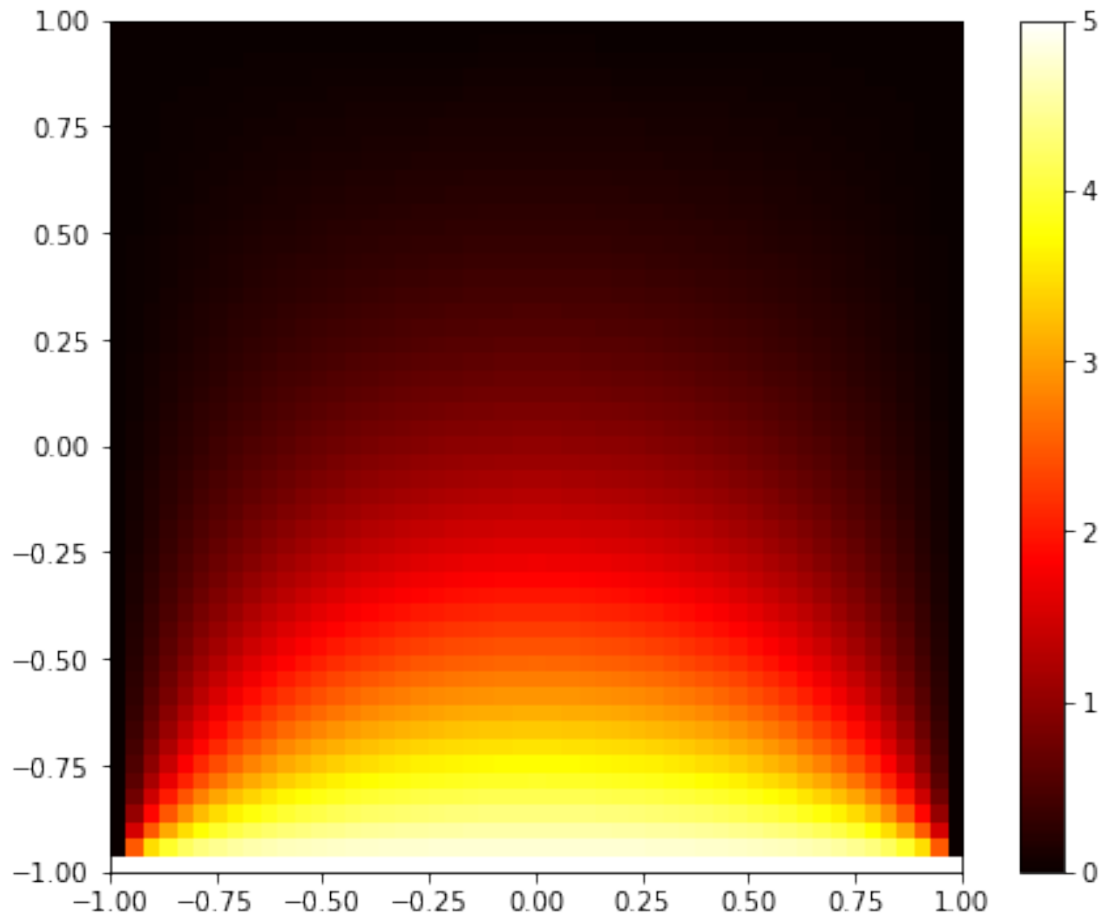
```

```

N: 51, M: 10001, C: 0.0625
t_star: 0.42409999999999696
t_actual: 0.424011387033
Relative error: 0.00020898723402139866
Time cost: 39.50638008117676 s

```

```
[ ]: <matplotlib.colorbar.Colorbar at 0x7fc9b9b02990>
```



0.3.2 By increasing the number of discretisation points demonstrate how many correct digits you can achieve.

For the Backward Euler method, there is no condition for C but the execution time increases without the help of GPU. Therefore, I choose smaller N and larger C to compute t^* .

```
[ ]: # set a list of N and C to research how the number of discretisation points  $N_{\square}$ 
      ↪ and time-steps  $M$  influence correct digits

t_actual = 0.424011387033
N_list = [11,51,81,101]
C_list = [1,0.5,0.25,0.125]

for N in N_list:

    print(f"N = {N}", end="\n")
    print(f"t_actual = {t_actual}")
```

```

t_exe = 0
table = PrettyTable(["C","M","t_star"],align = 'l')

for C in C_list:
    table_data = []

    A,_ = discretise_poisson(N)
    I = identity(N*N)

    u = np.zeros((N,N), dtype = np.float64)
    u[-1,:] = np.float32(5)

    dx = 2/(N-1)
    dt = C*dx**2
    M = round(1/dt + 1)
    table_data.append(C)
    table_data.append(M)

    t_star = 0
    with Timer() as t:
        while u[N//2,N//2] < 1:
            u = backward_euler(u,dt)
            t_star += dt

    t_exe += t.interval

    rel_err = abs(t_star-t_actual)/abs(t_actual)
    table_data.append(t_star)
    table.add_row(table_data)

print(table)
print(f"Time cost: {t_exe} s")
print("\n")

```

```

N = 11
t_actual = 0.424011387033
+-----+-----+-----+
| C      | M    | t_star                |
+-----+-----+-----+
| 1      | 26   | 0.44000000000000017  |
| 0.5    | 51   | 0.44000000000000003  |
| 0.25   | 101  | 0.43000000000000003  |
| 0.125  | 201  | 0.43000000000000003  |
+-----+-----+-----+
Time cost: 0.0996863842010498 s

```

```

N = 51
t_actual = 0.424011387033
+-----+-----+-----+
| C      | M      | t_star          |
+-----+-----+-----+
| 1      | 626    | 0.4255999999999983 |
| 0.5    | 1251   | 0.424800000000000395 |
| 0.25   | 2501   | 0.424400000000000805 |
| 0.125  | 5001   | 0.42419999999999859 |
+-----+-----+-----+
Time cost: 27.694674253463745 s

```

```

N = 81
t_actual = 0.424011387033
+-----+-----+-----+
| C      | M      | t_star          |
+-----+-----+-----+
| 1      | 1601   | 0.42437499999999373 |
| 0.5    | 3201   | 0.42437499999999234 |
| 0.25   | 6401   | 0.424062500000002255 |
| 0.125  | 12801  | 0.4240624999999826 |
+-----+-----+-----+
Time cost: 204.5022668838501 s

```

```

N = 101
t_actual = 0.424011387033
+-----+-----+-----+
| C      | M      | t_star          |
+-----+-----+-----+
| 1      | 2501   | 0.424400000000000805 |
| 0.5    | 5001   | 0.42419999999999859 |
| 0.25   | 10001  | 0.42409999999999696 |
| 0.125  | 20001  | 0.424049999999996146 |
+-----+-----+-----+
Time cost: 865.9391808509827 s

```

Similar with the conclusion for the Forward Euler method, when C is certain, the larger N and M is or the smaller Δt and h is, the more correct digits we can get.

0.3.3 Plot the convergence of the computed time * against the actual time.

I will use the same scheme as that for the Forward Euler method to plot in order to evaluate how the time steps and discretisation points influence t^* . Also, to save execution time, I have to choose smaller M and N .


```

[19]: # Plot the convergence graph against the number of discretisation points for a_
      ↪ certain number of time steps

M = 40001
N_list = [11,21,31,41,51,71]
err_list = []

for N in N_list:

    A,_ = discretise_poisson(N)
    I = identity(N*N)

    u = np.zeros((N,N), dtype = np.float64)
    u[-1,:] = np.float32(5)

    dx = 2/(N-1)
    dt = 1/(M-1)
    C = dt/dx**2

    t_star = 0
    while u[N//2,N//2] < 1:
        u = backward_euler(u,dt)
        t_star += dt

    rel_err = abs(t_star-t_actual)/abs(t_actual)
    err_list.append(rel_err)

plt.plot(N_list, err_list)
plt.xlabel("N: Number of discretisation points")
plt.ylabel("Relative Error")
plt.title("Fig.4 Implicit method: Convergence of t* to the actual time against_
      ↪ N")
plt.show()

plt.loglog(N_list, err_list)
plt.xlabel("N: Number of discretisation points")
plt.ylabel("Relative Error")
plt.title("Fig.5 Implicit method: Convergence of t* to the actual time against_
      ↪ N (loglog)")
plt.show()

```

Fig.4 Implicit method: Convergence of t^* to the actual time against N

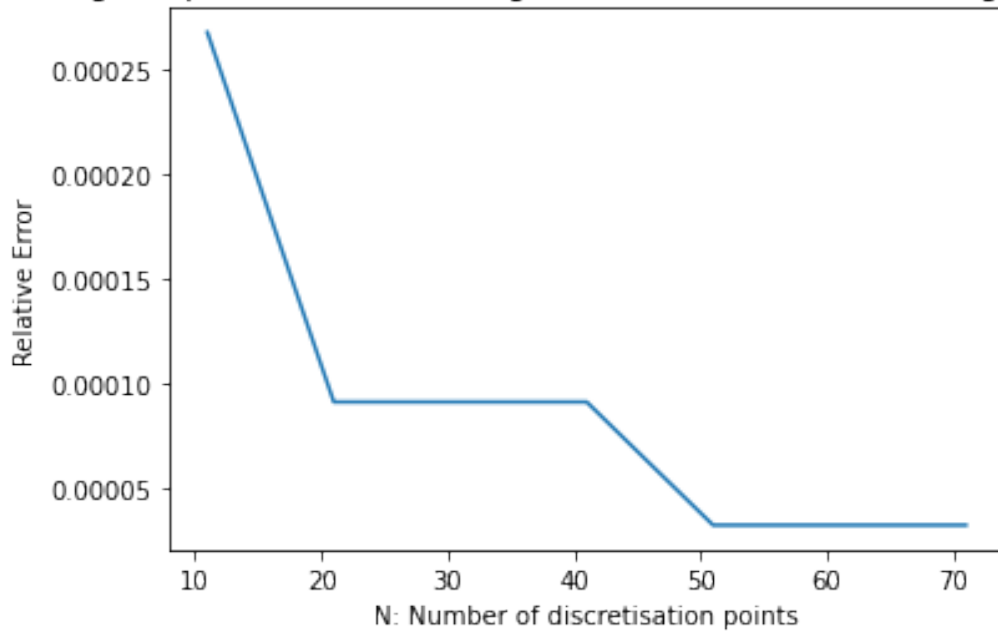
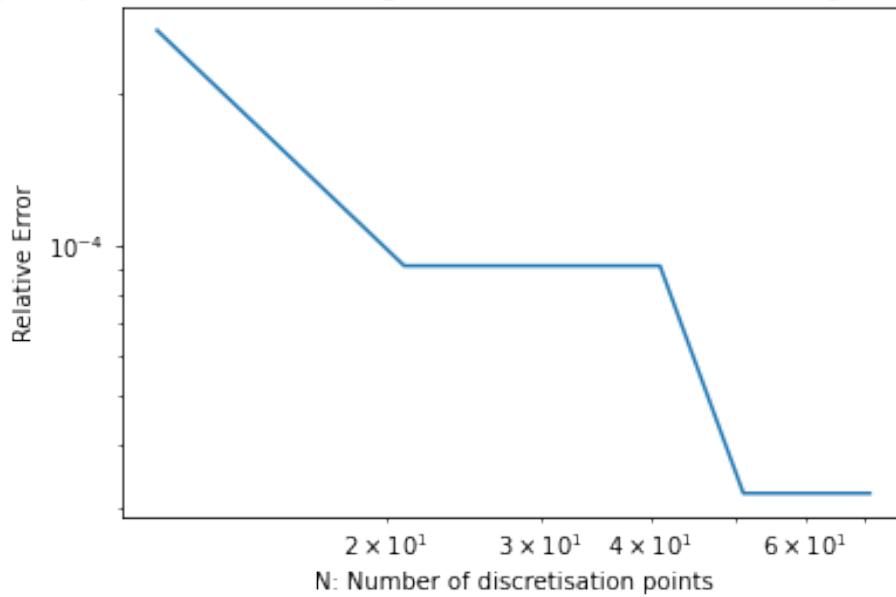


Fig.5 Implicit method: Convergence of t^* to the actual time against N (loglog)



It is obvious that when time steps are certain, t^* converges quadratically to t as discretisation points increase.

```

[ ]: # Plot the convergence graph against the number of time steps for a certain
      ↪ number of discretisation points

N = 101
M_list = [101,501,1001,3001,5001,8001,10001,15001]
err_list = []

for M in M_list:

    A,_ = discretise_poisson(N)
    I = identity(N*N)

    u = np.zeros((N,N), dtype = np.float64)
    u[-1,:] = np.float32(5)

    dx = 2/(N-1)
    dt = 1/(M-1)
    C = dt/dx**2

    t_star = 0
    while u[N//2,N//2] < 1:
        u = backward_euler(u,dt)
        t_star += dt

    rel_err = abs(t_star-t_actual)/abs(t_actual)
    err_list.append(rel_err)

plt.plot(M_list, err_list)
plt.xlabel("M: Number of time steps")
plt.ylabel("Relative Error")
plt.title("Fig.6 Implicit method: Convergence of t* to the actual time against_
      ↪ M")
plt.show()

plt.loglog(M_list, err_list)
plt.xlabel("M: Number of time steps")
plt.ylabel("Relative Error")
plt.title("Fig.7 Implicit method: Convergence of t* to the actual time against_
      ↪ M (loglog)")
plt.show()

```

Fig.6 Implicit method: Convergence of t^* to the actual time against M

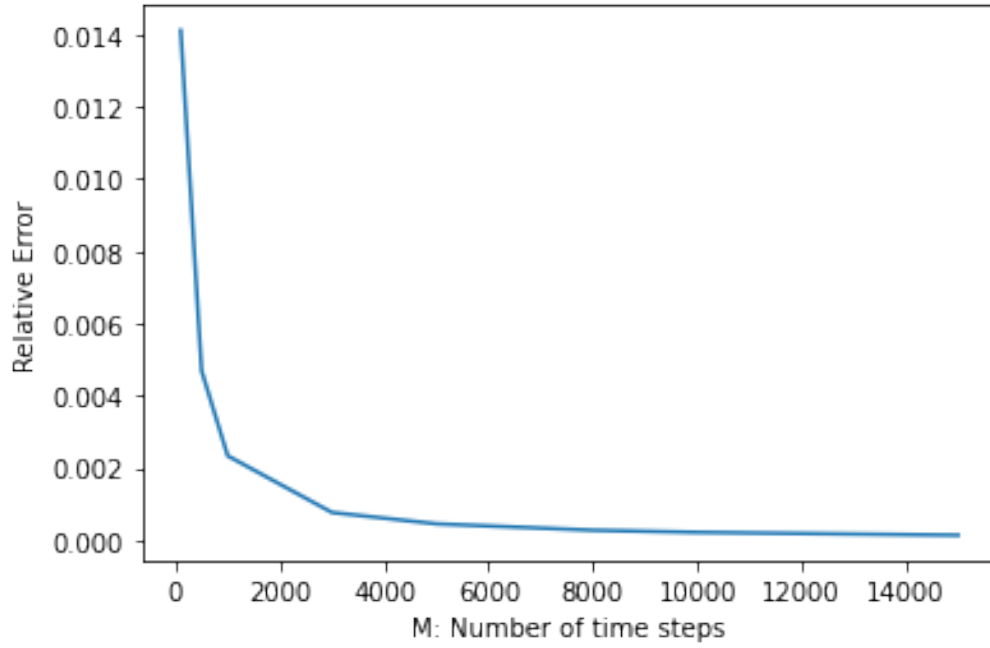
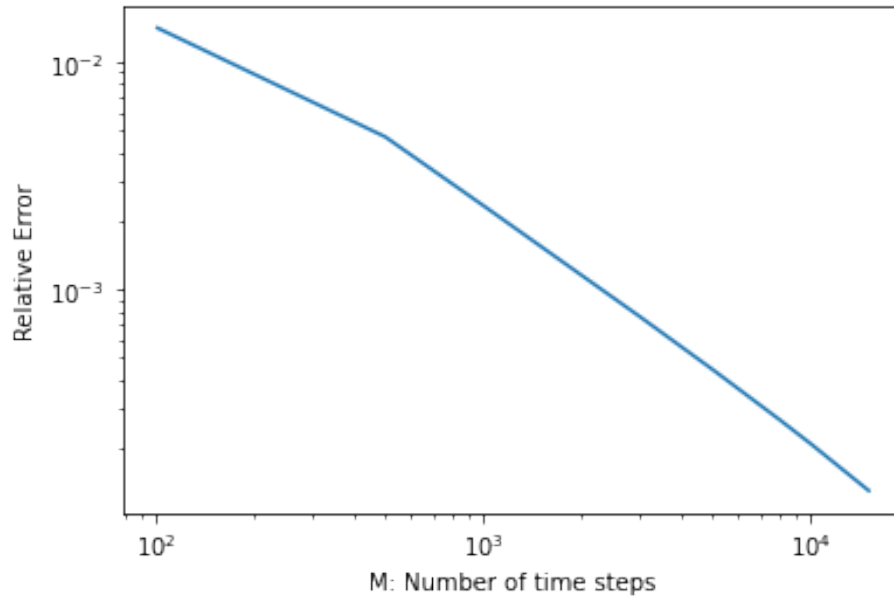


Fig.7 Implicit method: Convergence of t^* to the actual time against M (loglog)



Also, when N is certain, as M increases t^* quadratically converges to the actual t , which is similar to the Forward Euler method.